



Advanced Text Analysis for Business (IDS-566)

Lecture 2
Jan 26, 2018

Course Overview

- Instructor
 - Ehsan M. Ardehaly PhD, ehsan@uic.edu
 - Office hours: 4:45 - 5:45 pm F, BLC L270
 - Teacher assistant: 4:00 - 5:00 pm W, BLC L270
- Objectives:
 - Text mining
 - Applications for business decisions
 - Study of machine learning concepts
 - Design and implementation of text mining approaches

Course Overview

- Suggested text books:
 - Fundamentals of Predictive Text Mining (2nd Edition), Sholom M. Weiss, Nitin Indurkha, Tong Zhang, 2015
 - Mining Text Data, Charu C. Aggarwal and ChengXiang Zhai, Springer, 2012
 - [Mining of Massive Datasets](#), Jure Leskovec, Anand Rajaraman, Jeff Ullman
- Grading:
 - Final exam: 40%
 - 3 Assignments: 60% (3 x 20%)

Course Assignments

- Grade: 20%
- Loading textual data
- Building models
- Analysis
- Suggested programming language
 - Python 3
 - Scientific packages (e.g. scikit-learn)

Assignment policy

- Please read university regulations:
 - <https://grad.uic.edu/university-regulations>
- All assignment you turn must be done by you **alone**.
- The first violation will result in a failing grade for that assignment.
- The second will result in a failing grade for the course.
- **Late Submission Policy:**
 - 4% per hour
- **Grade dispute:**
 - Within 7 days of the receipt of the grade

Agenda

- Lexical analysis
 - Tokenization, Regular Expression
 - Unigram, Bigram, N-gram
 - Stop words
- Document to term matrix:
 - Vectorization
 - TF-IDF
- Sparse matrix:
 - LIL matrix, CSR matrix, CSC matrix
- Sentiment Analysis:
 - Lexicons

Lexical Analysis

- Text is unstructured
- Computer models need structure
- Solution:
 - Driving patterns from text
 - Creating structured data

Tokenization

- Sequence of characters → Sequence of terms
- Example:
 - Input: “Hello world!”
 - Word as a token:
 - Output: “hello”, “world”
 - Separate by space:
 - Output: “Hello”, “world!”

Tokenization

- Separating by special characters:
 - Space, tab, new-line, ...
 - What about Unicode?
 - Slow

Tokenization

- Separating by special characters:
 - Space, tab, new-line, ...
 - What about Unicode?
 - Slow
- Solution:
 - Using regular expression (regex)
 - Python: package 're'

Regular Expression

- A sequence of characters that define a search pattern:
 - Find operation
 - Find and replace operation
 - Find and split operation
- Example:
 - `[\n'()]` → Match space, new-line, single quote, parenthesis
 - `\w+` → Match multiple alphanumeric characters

Regular Expression use cases

- Tokenization to words:
 - `\w\w+`
- Recognizing URLs
 - `http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@.&+]|[*\(\),]|(?:%[0-9a-fA-F][0-9a-fA-F]))+`
- Recognizing hashtags (or mentions)
 - `(?<=^|(?<=[^a-zA-Z0-9-_\.]))#([A-Za-z]+[A-Za-z0-9]+)`

Unigram term

- Single word in sentence
- Without any space (or enter, tab, ...)
- Without special characters (!, ?, ...)
- Converting to lower case (optional)
- Sample: “this”, “is”, “i”

Bigram, n-gram

- Bigram: Contiguous sequence of two unigrams.
 - Example: “this is”, “is my”, “have you”, “text mining”
- N-gram: Contiguous sequence of N unigrams.
 - 3-gram: “text mining course”
 - 4-grams: “the president of us”
 - 5-grams: “the president of united states”

Tokenization to unigrams

- Input:
- “If I have enough money, I will go to Japan”
- Output:
- “if”, “i”, “have”, “enough”, “money”, “i”, “will”, “go”, “to”, “japan”

Tokenization to bigrams

- Input:
- “If I have enough money, I will go to Japan”
- Output:
- “if i”, “i have”, “have enough”, “enough money”, “money i”, “i will”, “will go”, “go to”, “to japan”

Unique unigram tokens:

- Input:
- “If I have enough money, I will go to Japan”
- Unique unigrams:
- “if”, “i”, “have”, “enough”, “money”, “will”, “go”, “to”, “japan”

Stop-words

- We often remove stop-words in tokenization.
- Some samples:
 - am, but, as, did, I, if, into, some, than, up, to, your, ...

Document

- Each sample of textual data:
 - Document 1: Shipment of gold damaged in a fire.
 - Document 2: Delivery of silver arrived in a silver truck.
 - Document 3: Shipment of gold arrived in a truck.
- Tokenization: unigram, lowercased, more than 1 letter.
 - arrived, damaged, delivery, fire, gold, in, of, shipment, silver, truck

Document to Term Matrix (DTM)

Each sample of textual data:

Document 1: Shipment of gold damaged in a fire.

Document 2: Delivery of silver arrived in a silver truck.

Document 3: Shipment of gold arrived in a truck.

Tokenization: unigram, lowercased, more than 1 letter.

arrived, damaged, delivery, fire, gold, in, of, shipment, silver, truck

Doc	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
D1	0	1	0	1	1	1	1	1	0	0
D2	1	0	1	0	0	1	1	0	2	1
D3	1	0	0	0	1	1	1	1	0	1

Count Vectorization vs. Binary Vectorization

Count vectorization (term frequency):

Doc	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
D1	0	1	0	1	1	1	1	1	0	0
D2	1	0	1	0	0	1	1	0	2	1
D3	1	0	0	0	1	1	1	1	0	1

Binary vectorization:

Doc	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
D1	0	1	0	1	1	1	1	1	0	0
D2	1	0	1	0	0	1	1	0	1	1
D3	1	0	0	0	1	1	1	1	0	1

Inverse Document Frequency

Term Frequency (TF):

Doc	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
D1	0	1	0	1	1	1	1	1	0	0
D2	1	0	1	0	0	1	1	0	2	1
D3	1	0	0	0	1	1	1	1	0	1

Inverse Document frequency (IDF) and log(IDF):

	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
IDF	3/2	3/1	3/1	3/1	3/2	3/3	3/3	3/2	3/1	3/2
log(IDF)	.41	1.1	1.1	1.1	.41	0	0	.41	1.1	.41

Smoothing IDF

$$idf = \log \frac{N}{DF} \quad \xrightarrow{\text{Smoothing}} \quad idf = 1 + \log \frac{N + 1}{DF + 1}$$

	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
IDF	3/2	3/1	3/1	3/1	3/2	3/3	3/3	3/2	3/1	3/2
log(IDF)	.41	1.1	1.1	1.1	.41	0	0	.41	1.1	.41
Smooth	1.29	1.69	1.69	1.69	1.29	1	1	1.29	1.69	1.29

TF-IDF

TF:

Doc	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
D1	0	1	0	1	1	1	1	1	0	0
D2	1	0	1	0	0	1	1	0	2	1
D3	1	0	0	0	1	1	1	1	0	1

IDF:

	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
Smooth	1.29	1.69	1.69	1.69	1.29	1	1	1.29	1.69	1.29

$$tf.idf = tf(1 + \log \frac{N + 1}{DF + 1})$$

Doc	arrived	damaged	delivery	fire	gold	in	of	shipment	silver	truck
D1	0	1.69	0	1.69	1.29	1	1	1.29	0	0
D2	1.29	0	1.69	0	0	1	1	0	3.39	1.29
D3	1.29	0	0	0	1.29	1	1	1.29	0	1.29

Normalizing TF-IDF

- Normalizing term vector
- Term vector: $[x_1 \quad x_2 \quad \dots \quad x_n]$
- L1: $[\frac{x_1}{\sum x_i} \quad \frac{x_2}{\sum x_i} \quad \dots \quad \frac{x_n}{\sum x_i}]$
- L2: $[\frac{x_1^2}{\sum x_i^2} \quad \frac{x_2^2}{\sum x_i^2} \quad \dots \quad \frac{x_n^2}{\sum x_i^2}]$

Zipf's law

- Empirical law:
 - There are many rare words.
 - There are a small number of extremely frequent words.
- Let f_i be the frequency of the i -th most common term:
- $f_i \propto \frac{1}{i}$
- Or
- $f_i \propto i^{-b}$

Storing Document to Term Matrix

- Normal array:
 - Dense
 - Most of cells are zero
 - To much memory
 - To many matrix operations

Storing Document to Term Matrix

- Normal array:
 - Dense
 - Most of cells are zero
 - To much memory
 - To many matrix operations
- Solution:
 - Sparse Matrix

Sparse Matrix

Lists of lists (LIL)

- Good for incremental matrix construction

Compressed Sparse Row (CSR)

- Good for row operations

Compressed Sparse Column (CSC)

- Good for column operations

LIL Matrix

- Good for incremental matrix construction

- $$\begin{bmatrix} 3 & 0 & 0 & 3 \\ 0 & 2 & 1 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}$$

- Data: $[[3, 3], [2, 1], [1]]$
- Rows: $[[0, 3], [1, 2], [1]]$

CSR Matrix

- Good for row operations

- $$\begin{bmatrix} 3 & 0 & 0 & 3 \\ 0 & 2 & 1 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}$$

- Data: $[3, 3, 2, 1, 1]$

- Indices: $[0, 3, 1, 2, 1]$

- Index pointer: $[0, 2, 4, 5]$



CSC Matrix

- Good for columns operations

- $$\begin{bmatrix} 3 & 0 & 0 & 3 \\ 0 & 2 & 1 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}$$

- Data: $[3, 2, 1, 1, 3]$

- Indices: $[0, 1, 2, 1, 0]$

- Index pointer: $[0, 1, 3, 4, 5]$

Example

- $X[2, 4] = 5$ ← Good for LIL matrix
- $Z = X.\text{dot}(Y)$
 - X: CSR matrix
 - Y: CSC matrix

Sentiment Analysis



7m

Congrats to @KingJames of the @cavs on joining the exclusive 30,000 Career Points Club!

#ThisIsWhyWePlay #AllForOne 



1.1K



1.5K

Sentiment Analysis

- Classification of documents/messages by sentiment
- Applications:
 - Recommendation systems
 - Opinion mining
 - Customer relation management
 - Market sentiment

Approaches

- Lexicons
 - Dictionary base
 - General purpose
- Machine learning
 - Learning from data
 - Customized to the application

Lexicon approach

- Dictionary of words with score
- Example: AFINN <http://neuro.imm.dtu.dk/wiki/AFINN>
 - crying -2
 - excited 3
 - hell -4
 - winner 4
 - smiled 2

AFINN lexicons

- Congrats to @KingJames of the @cavs on joining the exclusive 30,000 Career Points Club!
- Congrats: 2
- Exclusive: 2
- Average score: 2 > 0

Challenges

- I do not dislike cabin cruisers. (Negation handling)
- Disliking watercraft is not really my thing. (Negation, inverted word order)
- I'd really truly love going out in this weather! (Possibly sarcastic)

https://en.wikipedia.org/wiki/Sentiment_analysis#Application_in_recommender_system