

# Advanced Text Analysis for Business (IDS-566)

Lecture 5

Feb 23, 2018

# Course Overview

- Instructor
  - Ehsan M. Ardehaly PhD, [ehsan@uic.edu](mailto:ehsan@uic.edu)
  - Office hours: 4:45 - 5:45 pm F, BLC L270
  - Teacher assistant: 4:00 - 5:00 pm W, BLC L270
- Objectives:
  - Text mining
  - Applications for business decisions
  - Study of machine learning concepts
  - Design and implementation of text mining approaches

# Assignments-2

- Grade: 20%
- Sentiment analysis
- Due date: 2/25/2018
- Submission:
  - Notebook (code + analysis) → PDF
  - Word document with code as an appendix → PDF

# Agenda

## Artificial Neural Network

- Layers, activation, SGD

## Multi-Layer Perceptron:

- Word embedding

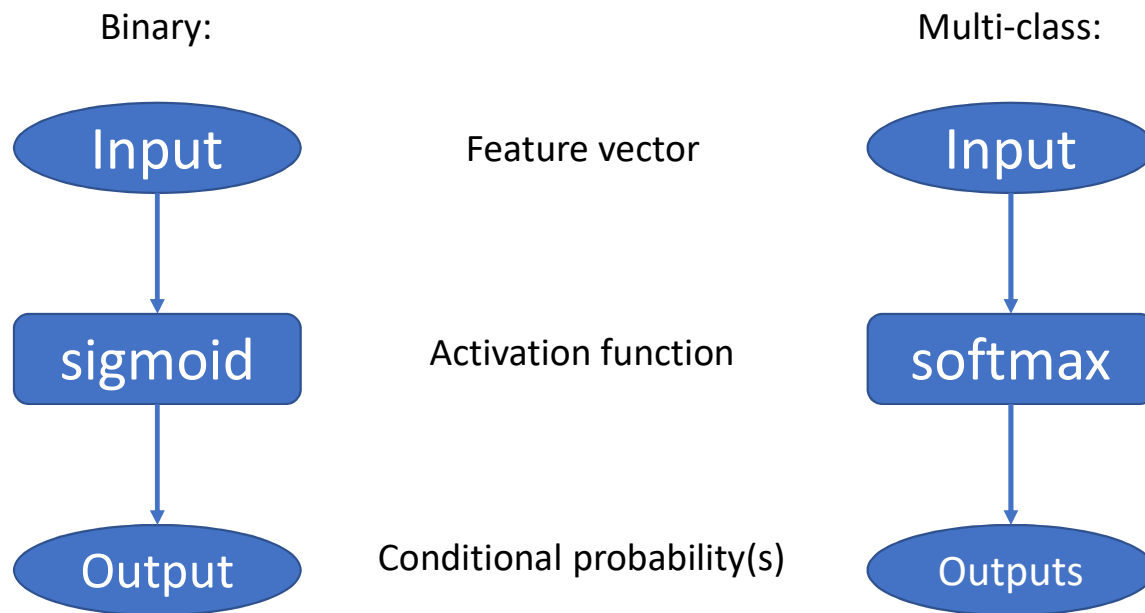
## Deep learning:

- Dropout, Batch normalization

## Examples:

- Keras

# Binary vs. Multi-class logistic regression



# Sigmoid vs. softmax

- Sigmoid (binary class):

- $P(y = 1|x) = \frac{1}{1+e^{-x.\theta}}$

- Softmax (multi-class):

- $P(y_i = k|x) = \frac{\exp(x.\theta^{(k)})}{\sum_{j=1}^m \exp(x.\theta^{(j)})}$

# Training logistic regression

- Creating the cost function:
  - Negative log likelihood
  - $J(\theta) = -l(\theta)$
- Find  $\theta$  which minimizes the cost function:
  - $\theta = \underset{\theta}{\operatorname{Argmin}} J(\theta)$

# Binary vs multi-class cost function

- Binary:
  - Negative log likelihood  $\rightarrow$  Binary cross-entropy
- Multi-class:
  - Negative log likelihood  $\rightarrow$  Categorical cross-entropy




# Gradient descent algorithm

- An iterative algorithm
- Finding the minimum of a cost function:
- Starts with a random initialization.
- Takes step to the negative of the gradient.

# Artificial Neural Network (ANN)


- Hidden layer:
  - Input: the output of last layer
  - Multiple cells
  - Apply multiple functions (per cell) → Activation function
  - Output: the next layer

# Hidden layer

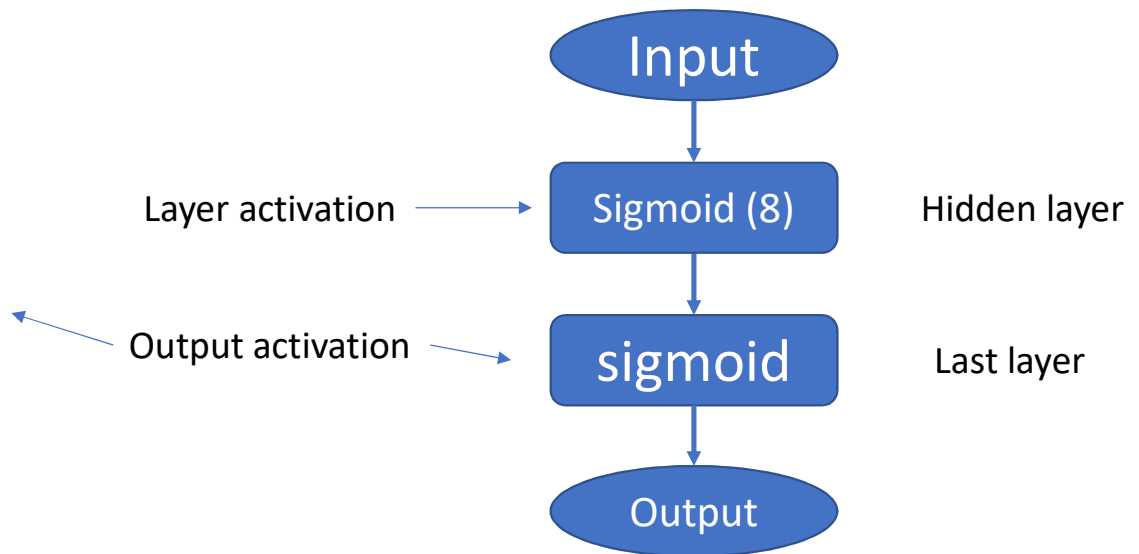


## 2 layers MLP (binary)

Binary Logistic Regression




MLP:

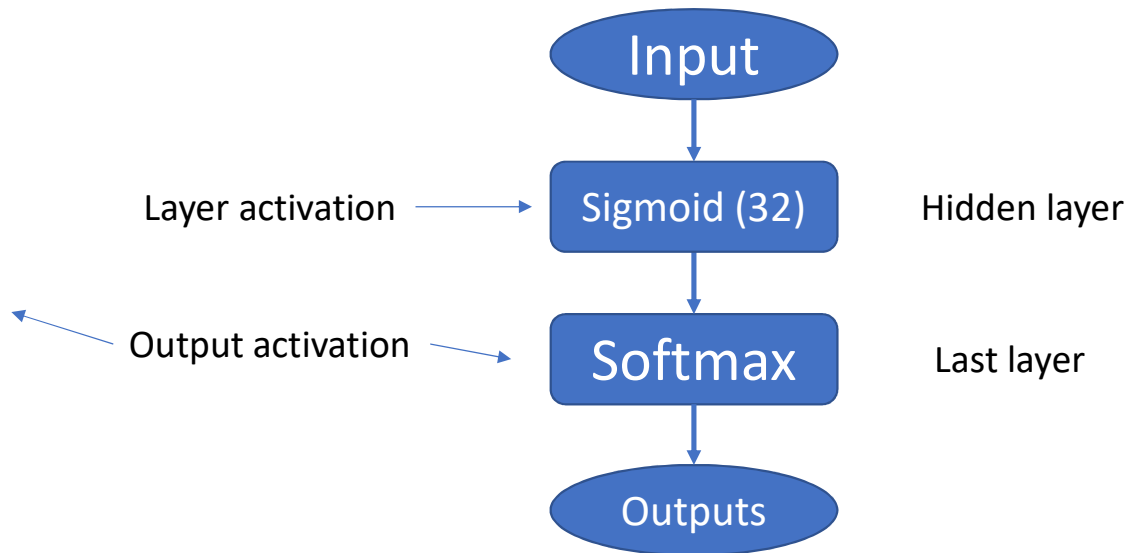


## 2 layers MLP (multi-class)

Binary Logistic Regression



MLP:




# Activation functions

- Middle layers:
  - Linear (not recommended)
  - Sigmoid
  - Relu (recommended)
  - Tanh
- Output layer:
  - Sigmoid (for binary)
  - Softmax (for multi-class)


# Rectifier linear unit (relu)

- $relu(x) = \max(0, x)$





# Tangent Hyperbolic

- $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- $\tanh(x) = 2\sigma(2x) - 1$






## 2 layer MLP (with relu)



## 3 layer MLP (with tanh)



# Optimization

## Cost function:

- Binary cross-entropy (binary)
- Categorical cross-entropy (multi-class)

## Batch gradient descent:

- Stochastic Gradient Descent (SGD)

# GD vs. SGD

## Gradient Descent:

Updating gradients to the entire data  
Good for simple models (e.g. logistic regression)  
All data must be fit in the memory.



## Stochastic Gradient Descent:

Dividing data to small batches

- Updating gradients for each batch.

Good for complex models (e.g. neural networks)  
Only a batch need to fit in the memory.

# Advanced optimizers (for batch)

- SGD with learning rate decay
  - Sensitive to the learning rate
- Adam:
  - Not sensitive to the learning rate
  - often converges very fast
- RMSprop:
  - Not sensitive to the learning rate

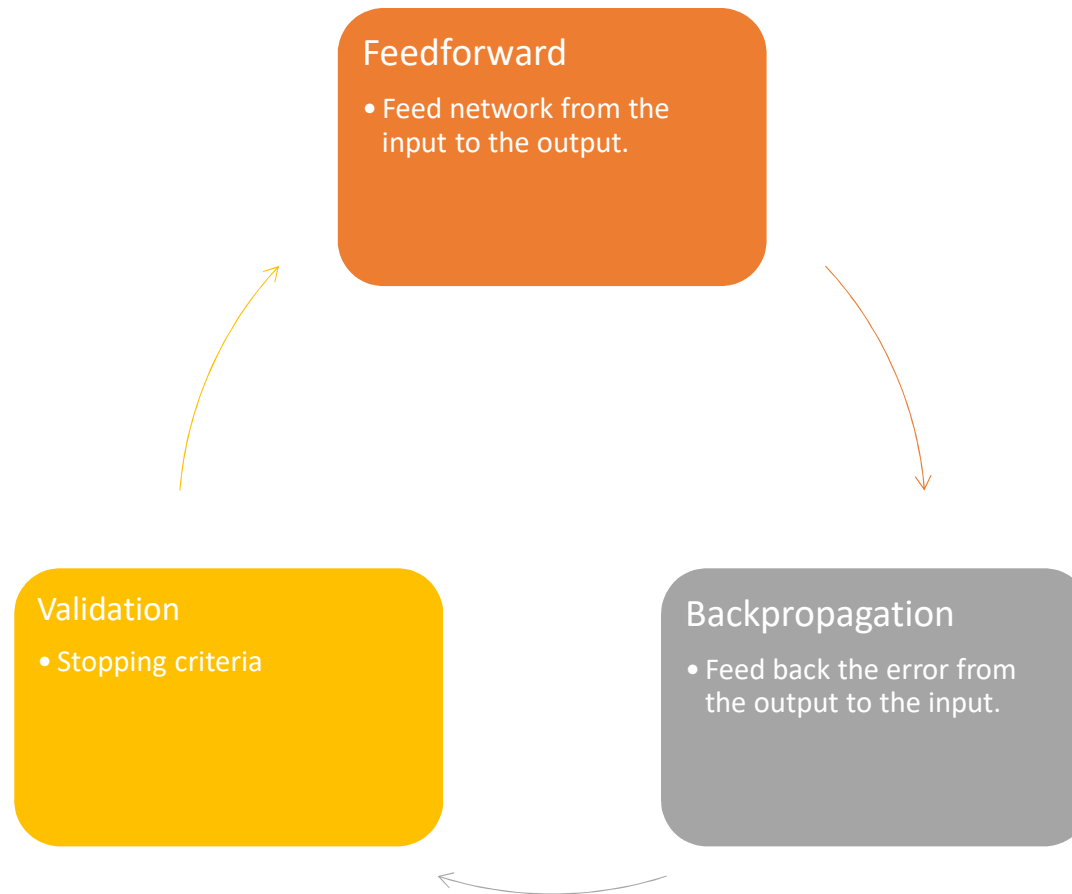
# Batch example - 1

- 100K instances with 2K features for 5 categories classification:
  - Shapes: X: 100000 x 2000, y: 100000
- Batch size: 1000
- Shuffle the data
- Create 100 batches:
  - Batch shapes: Shapes: X: 1000 x 2000, y: 1000

## Batch example - 2


- 100K instances with 2K features for 5 categories classification:
  - Shapes: X: (100000, 2000), y: (100000,)
- Batch size: 8000
- Shuffle the data
- Create 13 batches:
  - Batch shapes: X: (8000, 2000), y: (8000,)
  - Last batch shapes: X: (4000, 2000), y: (4000,)

# Training the network





# Feedforward




[https://en.wikipedia.org/wiki/Feedforward\\_neural\\_network](https://en.wikipedia.org/wiki/Feedforward_neural_network)

# Backpropagation


$$\frac{\partial E_{total}}{\partial w_1} = \frac{\partial E_{total}}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$

$$\downarrow$$

$$\frac{\partial E_{total}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial out_{h1}} + \frac{\partial E_{o2}}{\partial out_{h1}}$$



## Learning rate



<https://mattmazur.com/2015/03/17/a-step-by-step-backpropagation-example/>


# Deep learning

- A neural network with many layers
- Very complex networks
- Training is computation intensive

# Training on Nvidia GPUs


- Using Cuda cores
- Packages:
  - Tensorflow
  - Mxnet
  - Pytorch
  - Caffe

47X Higher Throughput than CPU Server on Deep Learning Inference



Workload: ResNet-50 | CPU: 1X Xeon E5-2690v4 @ 2.6GHz | GPU: add 1X NVIDIA® Tesla® P100 or V100

## Bias vs. Variance



# Regularization layers

- To reduce variance
- Different behavior in the training and the testing time
  - Dropout
  - Batch Normalization

# Dropout




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

Dropout: A Simple Way to Prevent Neural Networks from Overfitting, Srivastava, JMLR, 2014



# Batch Normalization

- Using a moving average to find the mean and the average of batch.
- Change the output of the layer to a zero mean vector with unit variance.
- Moving average updates only at the training time.

# Model weights

- Each layer may have multiple parameters.
- Dense layers have two parameters:
  - **Coefficient** vector per cell (Coefficient matrix).
  - **Bias** per cell (bias vector)
  - $l(z) = \tanh(z \cdot \omega + b)$

# Word embedding

- Language modeling
- Feature learning
- Words  $\rightarrow$  vector (in low dimension)

# Word2vect

- A MLP model
- Each word is mapped to the weights of the first layer.
  - First layer weight:  $\# \text{ words} \times \# \text{ cells}$
- We can use these vectors in different classification task.

# One hot encoding

- Multi-class classification
- Labels: [0, 1, 2, 1, 0]
- One hot encoding:

- $$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$